

Resume:-

1. LSTM, Bi-directional LSTM
 - why LSTM > RNNs
 - cell state vs hidden state
 - Difference b/w LSTM & Bi LSTM

2. NLP

- Transformer architecture
- self-attention
- Glove, ELMO, FastText!?

3. ** Information retrieval

- metrics:- MRR, NDCG
- Traditional methods

4. Classical ML

- Mathematics: Logistic, Linear Regression
SVM, PCA, SVD

- Linear algebra revision ✓

Chi-square test

- Hypothesis testing (linear regression)

Confidence interval

p-value
t-test, f-test etc.
non-parametric testing

Coding practice

Deep learning

- Optimisation

Local minima, saddle points, local maxima

→ Effect of sample size on p -value and confidence interval.

→ standard error

→ A/B testing case studies
think about business metrics } resources around this

→ Using statistics to evaluate choices

Boosting algorithm 1

Xgboost.

Bagging Random Forest

$$\begin{aligned}
 & VV^T = I \\
 & V^T = V^{-1} \\
 & (V^{-1})^{-1} = V \quad UU^T = I \quad UT = U^{-1} \\
 & (USV^T)^{-1} = (V^T)^{-1} S^{-1} U^{-1} \\
 & = \boxed{V S^{-1} U^T}
 \end{aligned}$$

Matrix Rank

→ number of linearly independent rows: row rank
 → number of linearly independent columns: column rank

$$\Rightarrow \text{row rank} = \text{column rank} = \text{Rank}$$

$$\boxed{\text{Rank}(m, n) \leq \min(m, n)}$$

For a square matrix if determinant is non-zero, square matrix is a full rank matrix.

SVD ^{general term} $A = USV^T$ ^{diagonal matrix} $S = \sigma_i (i=1, \dots, k) \text{ else } \sigma_i = 0$

^{compute SVD} $(m \times n) \quad m \times m \quad n \times n \quad n \times m$ $B = USV^T$ is the best k -rank approximation to A

U, V : orthogonal matrices $UU^T = I, VV^T = I$
 U, V : rotate, S : stretching (diagonal matrices)

$$\vec{y} = A\vec{x} = (USV^T)(\vec{x})$$

S : singular matrix with non-zero diagonal elements

→ if S has 0 element calculate pseudo-inverse

→ condition number = $\frac{\sigma_1 \text{ largest}}{\sigma_n \text{ smallest}}$

for all the points
 If the second derivative of function exists, it is convex if $\nabla^2 f(x) \geq 0$ i.e. whether all the eigenvalues of $\nabla^2 f(x)$ Hessian are non negative

→ Local minima is global minima for convex function
 → There can be multiple minima though

Positive definite matrix : Eigenvalues > 0

* If the Hessian is everywhere
positive semi-definite then function is
convex.

$$H = \begin{bmatrix} \frac{\partial^2 f}{\partial x_1^2} & \frac{\partial^2 f}{\partial x_1 \partial x_2} & \dots & \frac{\partial^2 f}{\partial x_1 \partial x_n} \\ \vdots & & & \\ \frac{\partial^2 f}{\partial x_n \partial x_1} & \frac{\partial^2 f}{\partial x_n \partial x_2} & \dots & \frac{\partial^2 f}{\partial x_n^2} \end{bmatrix}$$

1) All +ve eigenvalues \rightarrow local minimum
All -ve eigenvalues \rightarrow local maximum
mix \rightarrow saddle point

Momentum method

addition $\leftarrow V_t = \beta V_{t-1} + g_{t,t-1}$

state vector $x_t = x_{t-1} - \eta \cdot V_t$
 $\eta \rightarrow$ learning rate.

Information Theory

Entropy: $H(X) = - \sum_x P(x) \log P(x)$

Mutual information:

$$I(X, Y) = H(X) + H(Y) - H(X, Y)$$

How much does X inform Y

KL Divergence: \rightarrow model Amount of information lost
 $D_{KL}(P \parallel Q) = \sum_x P(x) \log \frac{P(x)}{Q(x)}$ when 'Q' is used to approximate 'P'.
 \downarrow
true distribution

* Maximizing log-likelihood of observing data X with respect to model parameters θ is equivalent to minimizing KL Divergence between the likelihood and the source distribution of the data.

Principle of maximum entropy

\rightarrow Max entropy distribution 'agrees with what is known, but expresses minimum uncertainty with respect to all other matters'.

Maximize $-\sum_{P(x)} P(x) \log P(x)$

constraints:- $\sum P(x) - 1 = 0$ $P(x)$ must integrate (sum) = 1

Max entropy distribution when:-

① over a finite discrete range $\{0, 1, \dots, N\}$
Uniform distribution

② continuous r.v. X with mean μ

$$\int_{-\infty}^{\infty} dx (x \cdot p(x)) - \mu = 0$$

\Rightarrow Exponential distribution

③ Maximum entropy distribution with a variance σ^2

\rightarrow variance constraint: $\int dx (x - \mu)^2 \cdot p(x) = \sigma^2$

Also implicit mean constraint

\Rightarrow Normal distribution

KL divergence measures the difference between two probability distributions over the same variable 'x'.

Logistic regression

MLE estimate:-

Bernoulli

$$P(Y=1 | X) = h_0(x) = P$$

$$P(Y=0 | X) = 1 - h_0(x) = 1 - P$$

$$\text{Likelihood} = (P)^y (1-P)^{1-y}$$

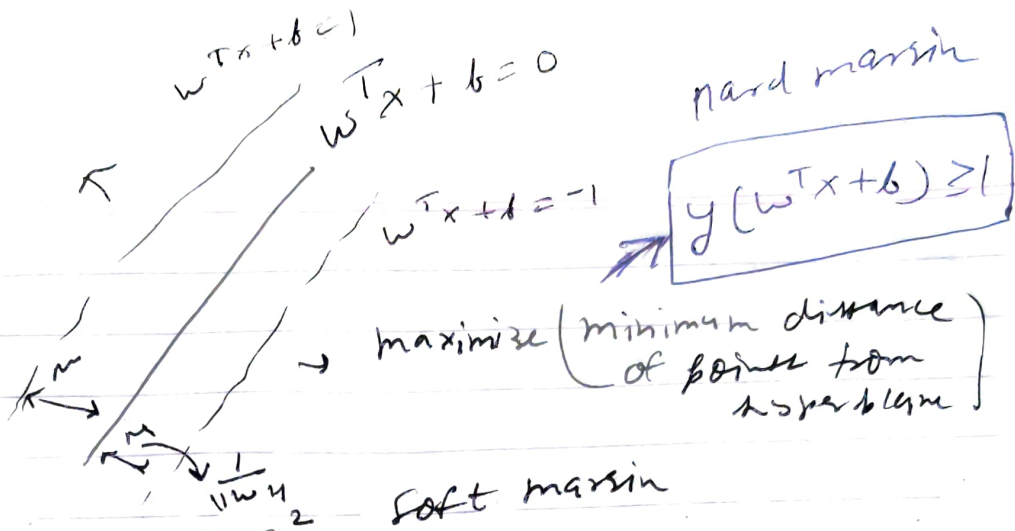
$$\frac{(W-k+2P)}{S} + 1 : \text{Output shape}$$

$$\Pr(|x - \mu| \geq k\sigma) \leq \frac{1}{k^2}$$

or equivalently

$$\Pr(|x - \mu| \geq k) \leq \frac{\sigma^2}{k^2}$$

SVM




$$\min. \frac{1}{2} \|w\|^2 + C \xi$$

such that $y_i (w_i x_i + b_i) \geq 1 - \xi_i$

Self-attention

sequence to sequence operation



Basic operation: $Y_i = \sum_j \left(\frac{w_{ij}}{\sum_j w_{ij}} \right) X_j$ weighted average of inputs

$w_{ij} = (X_i \cdot X_j) \rightarrow \text{softmax operation to normalize}$

$(X \cdot X^T) X$

keys, queries, value

$K = w_K^T X, Q = w_Q^T X, V = w_V^T X$

$\frac{\text{softmax}(K \cdot Q)}{\sqrt{\text{dim}}}(V)$ self-attention

SGD or GD

$$\theta = \theta - \alpha \nabla L_{\theta}$$

Ad Momentum

$$v = \beta v_{t-1} + \nabla L_{\theta}$$

$$\theta = \theta - \alpha v$$

Adagrad

$$r = r + \nabla L_{\theta} \odot \nabla L_{\theta}$$

$$\theta = \theta - \frac{\alpha}{\epsilon + \sqrt{r}} \nabla L_{\theta} \rightarrow g: \text{gradients}$$

RMS Prop with momentum

$$r = (1 - \rho) \nabla L_{\theta} \odot \nabla L_{\theta} + \rho(r) \quad \text{decay factor} \quad \text{exponential average}$$

$$v = \beta v + \alpha \frac{1}{\sqrt{r}} \odot g$$

$$\theta = \theta - \alpha v$$

decay factors

Adam

$$s = \rho_1 s + (1 - \rho_1) g$$

$$r = \rho_2 r + (1 - \rho_2) g \odot g$$

$$s = \frac{s}{1 - \rho_1}$$

$$r = \frac{r}{1 - \rho_2}$$

$$\theta = \theta - \alpha \frac{s}{r}$$

SVD relation to Eigen-decomposition

SVD:

$$A = U \Sigma V^T$$

Eigen decomposition

$$A = X \Lambda X^T \quad *$$

A needs to be symmetric
 U, V, X are orthonormal
 Λ, Σ are diagonal

$$AA^T = (U \Sigma V^T)(U \Sigma V^T)^T = (U \Sigma V^T)(V \Sigma^T U^T)$$

$$= U \Sigma \Sigma^T U^T$$

$$\begin{matrix} \downarrow & \downarrow & \downarrow \\ X & \Lambda & X^T \end{matrix}$$

$$A^T A = \begin{matrix} & \uparrow & \uparrow & \uparrow \\ V & \Sigma & \Sigma^T & V^T \\ \uparrow & \uparrow & \uparrow \\ X & \Lambda & X^T \end{matrix}$$

Shows how to SVD using Eigenvalue decomposition

$$\lambda_i = \sigma_i^2$$

SVD:-

1. Optimal low-rank approximation
2. Interpretability problem
3. Lack of sparsity

LSTM

3 gates: Input, forget and Output

$$I = \sigma(\quad)$$

$$f = \sigma(\quad)$$

$$O = \sigma(\quad)$$

Additional memory cell C_t : using H_{t-1}

candidate $\tilde{C}_t = \tanh[X_t, H_{t-1}, W]$ ✓

$C_t = \underbrace{\tilde{C}_t}_{\text{memory cell}} \rightarrow \text{forget from the previous cell state}$

$$C_t = F_t \odot C_{t-1} + I_t \odot \tilde{C}_t \rightarrow \text{read from the new cell state}$$

$$\boxed{H_t} = O_t \odot \tanh(\underline{C_t})$$

Gated Recurrent Unit :-

RNNs :-

$$H_t = f(W_h, X_t, H_{t-1}) \quad \checkmark$$


$$O_t = g(H_t, W_o) \quad \checkmark$$

GRUs introduce two designed (gates) :-

1. Reset gate
2. Update gate

Hidden state

H_{t-1}



Sigmoid gates

$$R_t = \sigma(X_t w_{xr} + H_{t-1} w_{hr} + b_r) \quad \checkmark$$

$$z_t = \sigma(X_t w_{xz} + H_{t-1} w_{hz} + b_z) \quad \checkmark$$

Input X_t

Reset gate controls how much previous state we might want to remember

Update gate controls how much of the new state is just a copy of the old state

Candidate hidden states :-

$$\tilde{H}_t = \tanh(X_t w_{xh} + (R_t \odot H_{t-1}) w_{rh} + b_h)$$

element wise multiplication

$$H_t = z_t \odot H_{t-1} + (1 - z_t) \odot \tilde{H}_t$$

helps with vanishing gradient problem in RNNs and better captures the dependencies for sequences with large time step distances